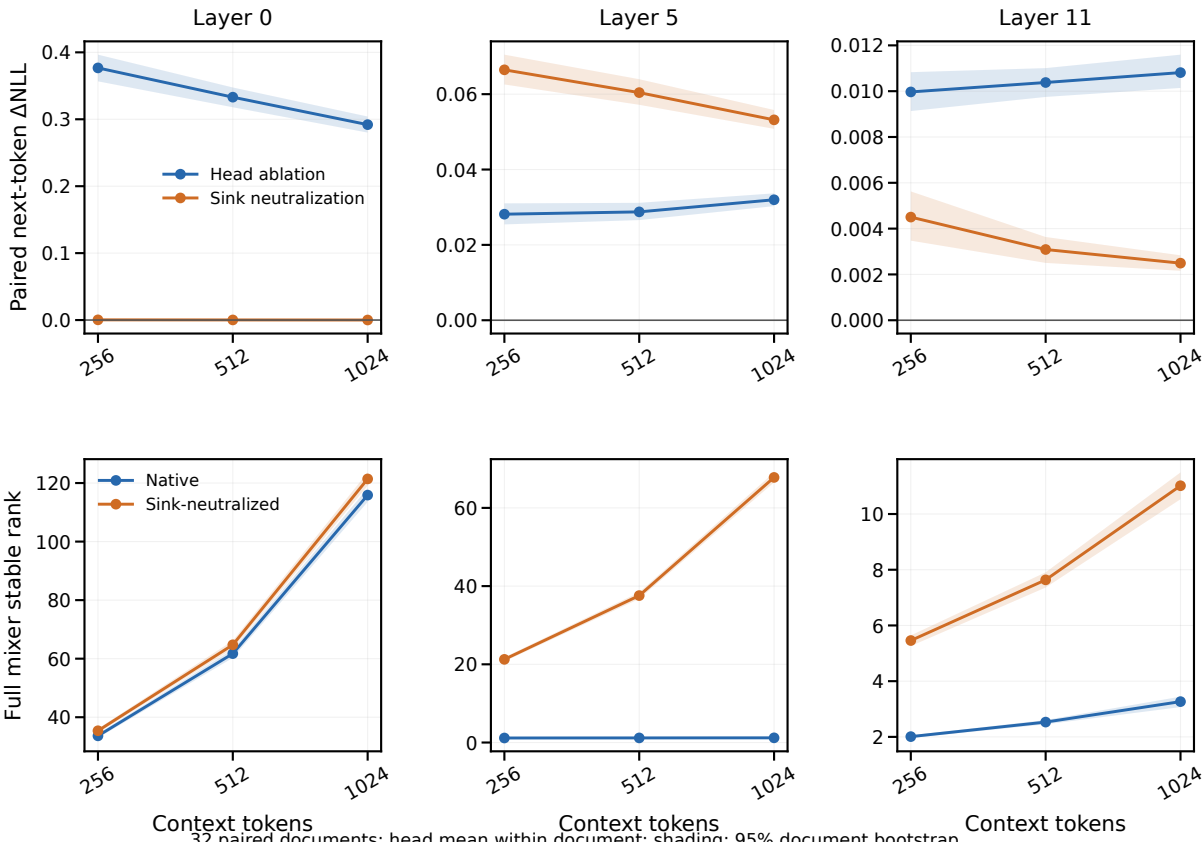
GPT-2: separately applied single-head interventions
Layer 0
Layer 5
Layer 11
Paired next-token ΔNLL
Head ablation
Sink neutralization
Full mixer stable rank
Native
Sink-neutralized
Context tokens
Context tokens
Context tokens
32 paired documents; head mean within document; shading: 95% document bootstrap.